COS30019 – Introduction to AI

Assignment 1 – Research into AI Ethics & Responsible AI

Due: 11:59pm Sunday 12th April 2026

Student name: Frederick Bartholomew Dao Xiang SII

Student ID: 104397146

Table of Contents

1. Executive Summary 3
2. Introduction 3
3. Findings 4
3.2 Business Incentives vs Ethical AI 4
3.2 Implementation Challenges of Responsible AI in Organisations 6
3.3 Why Companies Still Engage in Unethical AI Practices 7
3.4 Corporate Governance and Accountability in AI 8
3.5 AI in Government and Defence Contracts in the US 10
3.6 Effectiveness of Current AI Ethics Frameworks 11
4. Conclusion and Recommendations 12
5. Acknowledgements/Resources 13
References 14

1. Executive Summary

Artificial Intelligence (AI) has shown to be vital in staying competitive in the current modern business world. When companies compete to deploy AI systems quickly, it raises questions about the balance between ethics and economic incentives. Businesses tend to prioritize market advantage and position, speed, and profitability, while ethics focus on transparency, accountability, fairness and virtues. This research report focuses on the tension between ethics and businesses in AI.

This report finds that the tension between business incentives and ethical AI is driven by structural factors within competitive markets, where profitability, efficiency, and rapid innovation often outweigh ethical considerations. While ethical frameworks such as Australia's AI Ethics Principles provide valuable guidance, their voluntary and hard-to-enforce nature limits their effectiveness in practice. Organizations face significant challenges in implementing responsible AI due to ambiguous principles, technical limitations, and resource constraints, which contribute to inconsistent adoption of ethical practices. Furthermore, companies may engage in unethical AI practices as a rational response to competitive pressure and weak regulatory enforcement, while corporate governance frameworks are often insufficient due to issues such as self-regulation and ethics washing. These challenges are further intensified in high-stakes contexts such as government and defence contracts, where accountability and human oversight become more complex. Overall, the report concludes that achieving responsible AI requires stronger regulatory frameworks, improved organizational accountability, and better alignment between business incentives and ethical outcomes.

2. Introduction

Artificial Intelligence has been increasingly implemented into businesses and corporations over recent years, creating big economic changes and opportunities that understandably raises a lot of ethical concerns. As Hagendorff pointed out, there is significant links between corporations and the use of AI by the evidence that there has been an increasing number of research conducted for the use of AI towards the interests of the corporations, who are funding said research (Hagendorff 2020; Maslej et al. 2025). Hagendorff also states in the same paper, that businesses tend to prioritize profits and efficiency, and the topic of ethics can be omitted as he tends that "skipping ethical considerations is equivalent to the path of least resistance" in business contexts. Governments and research institutions have therefore begun developing ethical guidelines and frameworks to ensure that AI systems are created responsibly. One example, Australia's AI Ethics Principles, outline key values such as fairness, reliability, privacy protection, and human oversight in AI systems (*Australia's AI Ethics Principles* 2025). However, despite the creation of these frameworks, there has been tensions that commonly occur between ethical considerations and commercial incentives for the use of AI in businesses.

One major driving force for the adoption of AI in businesses is due to competitive pressures that prioritize profitability, efficiency and speed of innovation to capture a market as quickly as possible. These priorities may conflict with ethical principles that require careful risk assessment, transparency, and accountability. Wright & Schultz suggest that organizations face difficult trade-offs when implementing responsible AI practices because ethical considerations may increase development time or increase operational costs (Wright & Schultz 2018).

Ethical challenges associated with AI adoption are further exacerbated by issues of governance and corporate social responsibility. Camilleri supposes that effective AI ethics governance requires organizations to implement oversight mechanism, ethical guidelines, and accountability structures to ensure responsible decision making (Camilleri 2024). However, in cases where governance of these frameworks are weak, businesses may still engage in unethical

practices wherever the commercial incentives outweigh the ethical considerations (Méndez-Suárez et al. 2023). Additionally, implementing AI ethically within business operations often presents itself as a practical challenge; biases in training data, lack of transparency in algorithmic decision-making, and difficulties in ensuring compliance with ethical standards are commonplace (Purwanto et al. 2024).

This report investigates the tension between ethics and business incentives in the development and deployment of AI systems. A compilation of academia, ethical frameworks, and policy guidelines allows me to examine the key factors that influence ethical decision-making in organizations that are adopting AI technologies. The research focuses on how businesses balance innovation and profitability with ethics and evaluates whether existing ethical principles and governance frameworks are sufficient to encourage responsible AI practices.

3. Findings

3.2 Business Incentives vs Ethical AI

The rapid adoption of AI in business environments has created a fundamental tension between commercial objectives and ethical responsibilities. On one hand, organizations are driven by incentives such as profits, operational efficiency, innovation speed, and competitive advantage. On the other hand, ethical AI frameworks emphasize principles such as fairness, transparency, accountability, privacy protection, and the avoidance of harm. While both sets of priorities are important, they often conflict in practice, especially when ethical considerations impose additional costs, delays, or operational constraints on AI development and deployment.

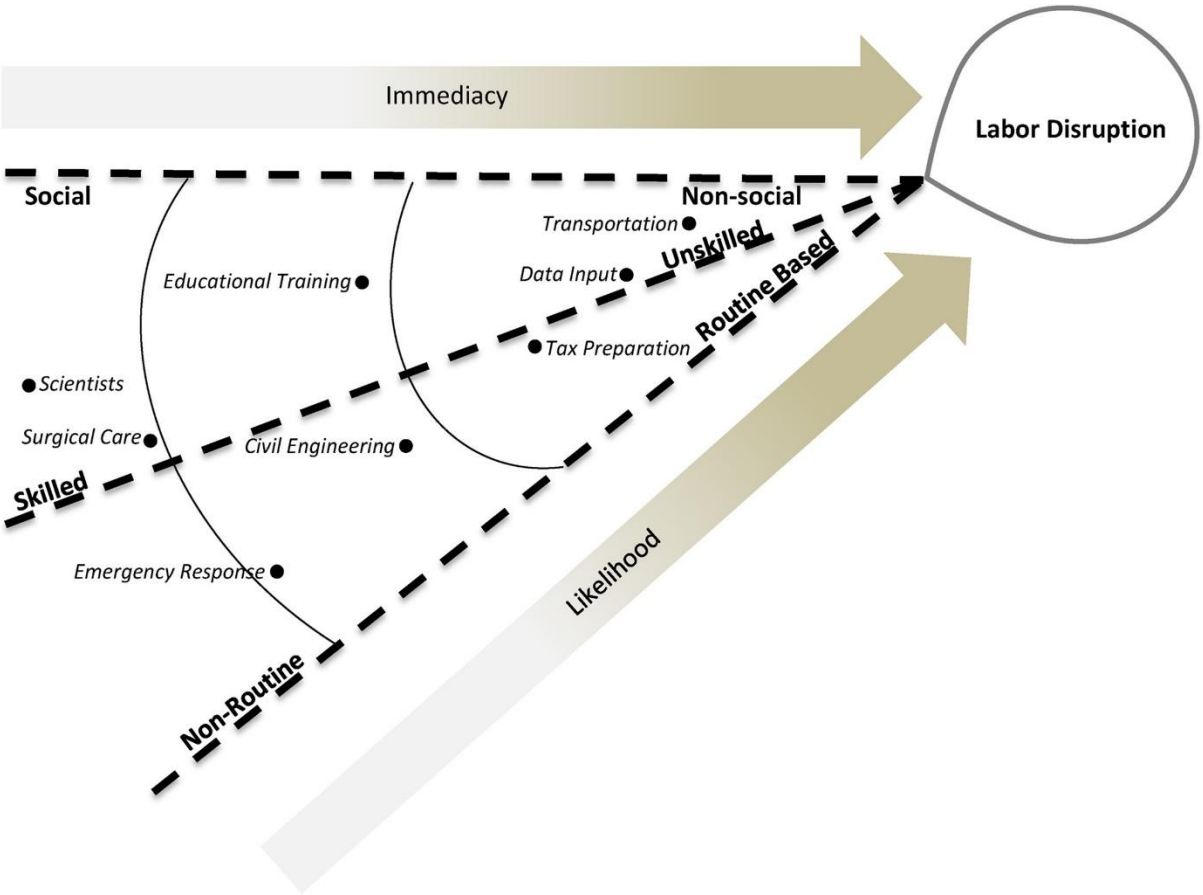

FIGURE 1 FACTORS INFLUENCING THE LIKELIHOOD AND IMMEDIACY OF AUTOMATION-BASED LABOR DISRUPTIONS (WRIGHT & SCHULTZ 2018)

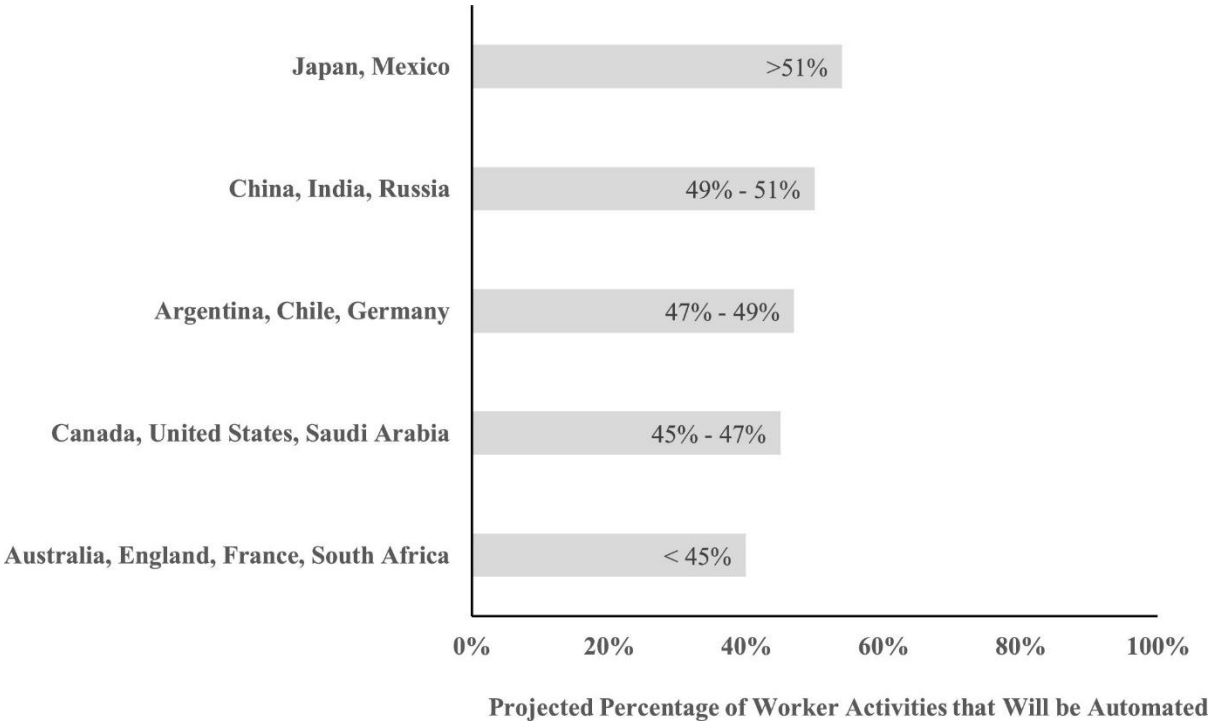

FIGURE 2 PROJECTED PERCENTAGE OF WORKER ACTIVITIES THAT WILL BE AUTOMATED BY REGION (WRIGHT & SCHULTZ 2018)

Businesses operate in highly competitive markets where the speed of innovation is a critical determinant of success. The deployment of AI systems can provide significant advantages, including cost reduction, improved decision-making, and enhanced customer targeting. As a result, companies may prioritize rapid development and implementation over thorough ethical evaluations. Research suggests that ethical AI practices, such as bias auditing, explainability measures, and governance oversight, can require a lot of time and financial investment, which may discourage organizations and businesses from fully adopting them (Wright & Schultz 2018). This creates a situation where ethical considerations are treated as secondary to business performance, particularly in industries where the first ones who innovate have a major advantage in their respective sectors. Wright and Schultz also propose that middle-skilled jobs are being replaced by automation, but low and high-skilled workers are being employed more, suggesting a polarization in employment. Figure 1 and Figure 2 shows their proposed cause, and projection.

In response to these challenges, governments and institutions have introduced ethical frameworks to guide responsible AI development. For example, Australia's AI Ethics Principles promote values such as fairness, transparency, accountability, and human-centred design (*Australia's AI Ethics Principles* 2025). However, critics argue that such frame works are often too abstract and lack accountability, limiting their effectiveness in real-world business contexts. Hagendorff highlights that many AI ethics guidelines share similar high-level principles but fail to provide concrete mechanisms for implementation or accountability. (Hagendorff 2020) As a result, organizations may adopt these principles in theory while failing to meaningfully integrate them into practice.

This gap between ethical ideals and business realities is further exacerbated by the voluntary nature of most AI ethics frameworks. Without strong regulatory enforcement, companies may selectively implement ethical practices or engage in "ethics washing", where they publicly promote responsible AI commitments without making substantive changes to their operations. In such cases, ethical frameworks function more as reputational tools rather than binding constraints on business behaviour.

Ultimately, the tension between business incentives and ethical AI reflects a broader structural challenge in the governance of emerging technologies. While ethical principles provide an important foundation for responsible innovation, they must contend with powerful economic forces that prioritize efficiency, growth, and market dominance. Addressing this tension requires note only clearer and more enforceable guidelines but also a shift in how organizations integrate ethical considerations into their core strategic decision-making processes.

3.2 Implementation Challenges of Responsible AI in Organisations

While ethical AI principles provide an important foundation for responsible development, organizations often face significant challenges when attempting to translate these principles into practice. The gap between high level ethical guidelines and real-world implementation is widely recognized in both academic and industry literature. Although frameworks such as Australia's AI Ethics Principles promote values like fairness, transparency, and accountability, applying these concepts within complex AI systems is far from straightforward.

One major challenge lies in the ambiguity of ethical principles themselves. Concepts such as "fairness" and "transparency" are inherently difficult to define in technical terms and can vary depending on context. For example, fairness in AI systems may involve equal outcomes, equal

opportunities, or the removal of bias, each of which can require different, and sometimes conflicting, technical approaches. This lack of standardization makes it difficult for organizations to determine which ethical criteria to prioritize and how to measure compliance effectively. As noted by Hagendorff, many AI ethics guidelines remain abstract and fail to provide concrete operational guidance, leaving organizations to interpret and implement principles on their own. (Hagendorff 2020)

In addition to conceptual challenges, organizations also face practical and resource related barriers. Implementing responsible AI practices often requires significant investment in specialized expertise, governance structures, and technical tools. Activities such as bias detection, model auditing, explainability enhancement, and ongoing monitoring can be costly and time consuming. For many, especially small and medium-sized organizations, these requirements may be difficult to justify in the face of competing business priorities. Purwanto and Fauzan suggest that without clear regulatory incentives or market pressure, companies may deprioritize ethical considerations in favour of efficiency and cost reduction. (Purwanto et al. 2024)

Technical limitations further complicate the implementation of ethical AI. Many advanced AI systems, particularly those based on deep learning and generative AI, operate as “black boxes”, making it difficult to fully understand or explain their decision-making processes. This issue becomes even more pronounced in the context of emerging agentic AI systems, which are capable of autonomous goal-based behaviour and decision-making. In high stakes domains such as national security and defence, these systems introduce additional layers of complexity, including unpredictability, scalability risks, and challenges in maintaining human oversight. Joshi highlights that ensuring accountability and reliability in such system requires not only technical safeguards but also comprehensive governance frameworks and policy interventions. These requirements further increase the difficulty of implementing ethical AI, particularly in environments where rapid deployment and strategic advantage are prioritized.

Organizational culture and governance also play a critical rose in shaping how ethical AI is implemented. Effective AI governance requires clear policies, oversight mechanisms and accountability structures to ensure that ethical considerations are integrated into decision-making processes. However, studies have shown that many organizations lack mature governance frameworks or fail to prioritize ethical oversight at a strategic level. (Camilleri 2024) In some cases, ethical responsibility may be diffused across teams, leading to unclear accountability and inconsistent implementation of ethical standards.

Overall, these challenges highlight that the implementation of responsible AI is not merely a technical problem but a complex organizational issue involving trade-offs between ethical ideals and practical constraints. Without clearer standards, stronger incentives, and more robust governance mechanisms, organizations are likely to continues struggling to operationalize ethical AI principles effectively.

3.3 Why Companies Still Engage in Unethical AI Practices

Despite the growing availability of ethical guidelines and governance frameworks, many organizations continue to engage in unethical or questionable ai practices. This persistence is not merely the result of oversight or technical limitations, but often reflects deliberate trade-offs shaped by economic incentives, competitive pressures, and regulatory gaps. Understanding why such behaviour occurs is essential to evaluating the effectiveness of current approaches to responsible AI.

A primary factor driving unethical AI practices is the prioritization of profit and competitive advantage. In many industries, AI systems provide significant opportunities for cost reduction, revenue generation, and market differentiation. As a result, organizations may be incentivized to deploy AI systems quickly, even when ethical risks have not been fully addressed. Méndez-Suárez argues that companies may knowingly adopt prohibited or ethically questionable AI practices when the expected benefits outweigh potential risks, especially in environments where enforcement mechanisms are weak or inconsistent (Méndez-Suárez et al. 2023). In such cases, unethical behaviour can be understood as a rational response to market conditions rather than an anomaly.

Competitive pressure further reinforces this dynamic. Firms operating in fast-moving technological sectors often compete to be the first to innovate, in which delaying deployment to make sure ethical guidelines are followed may result in lost market share or a disadvantage in the market. This creates a collective problem, as while ethical compliance may benefit society, individual firms have incentives to cut corners if competitors if competitors are perceived to be doing the same. Consequently, even organizations that recognize the importance of ethical AI may feel pressured to compromise to remain competitive.

Another key issue is the lack of strong regulatory oversight and enforcement. Many existing AI ethics frameworks are simply voluntary and lack legal consequences. This regulatory gap reduces the perceived cost of unethical behaviour, especially when the chance of being caught or punishment is low. In such situations, ethical guidelines may function more as “virtue symbols” rather than enforceable standards, enabling practices that prioritize business outcomes, over ethical considerations.

Organizational factors also contribute to unethical practices. Decision making within companies is often split across technical, managerial, and executive levels, which can mask accountability for ethical outcomes. Additionally, ethical considerations may be deprioritized if they are not integrated into core business strategies or performance metrics. Without clear incentives, employees and/or teams may focus primarily on technical performance and business objectives over broader ethics.

Overall, the continued presence of unethical AI practices highlights a structural misalignment between ethical expectations and business realities. As long as organizations operate in environments where economic incentives favour rapid innovation and weak oversight, ethical compromises have a high chance of persisting. Addressing this issue requires not only stronger regulatory frameworks, but also fundamental changes in how organizations align ethical responsibility with business outcomes.

3.4 Corporate Governance and Accountability in AI

In response to the ethical challenges posed by artificial intelligence, organizations have increasingly turned to corporate governance mechanisms to promote responsible AI development and deployment. These mechanisms include internal ethics committees, AI governance frameworks, risk management processes, and corporate social responsibility initiatives. In principle, such structures are designed to ensure that ethical considerations are integrated into organizational decision making and that accountability is maintained throughout the lifecycle of AI systems. However, the effectiveness of these governance approaches remains contested.

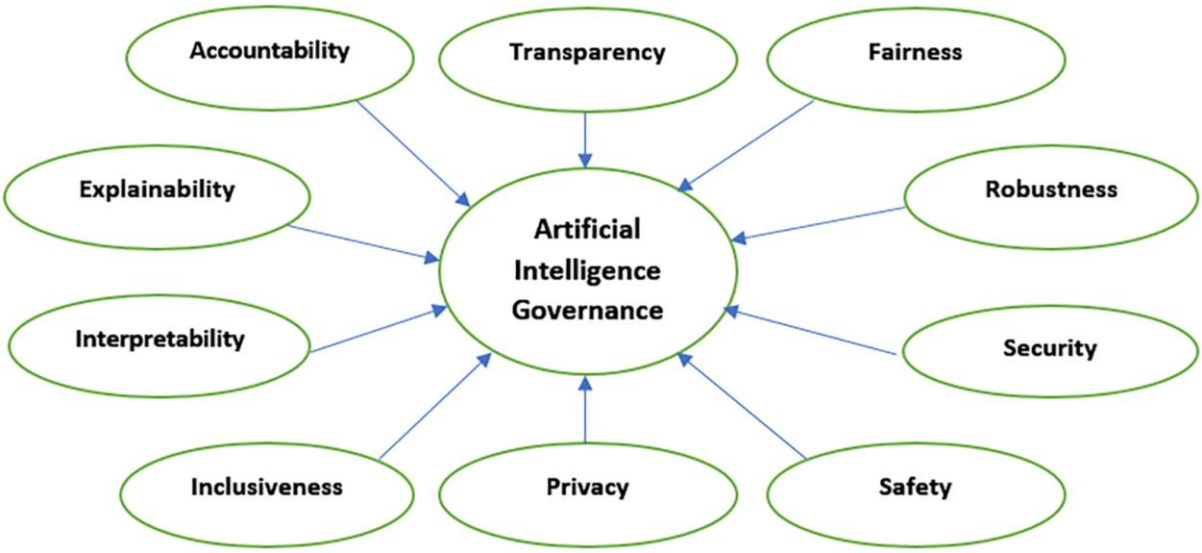

FIGURE 3 A RESPONSIBLE ARTIFICIAL INTELLIGENCE GOVERNANCE FRAMEWORK.(CAMILLERI 2024)

A key component of AI governance is the establishment of internal oversight structures, such as ethics boards, which are tasked with evaluating the potential risks and societal impacts of AI systems. These bodies aim to provide guidance on issues such as bias, transparency, and fairness, while also ensuring compliance with relevant regulations and ethical standards. Camilleri emphasizes that effective governance requires not only formal policies but also a broader commitment to corporate social responsibility, where organizations recognize their obligations to shareholders beyond their shareholders (Camilleri 2024). Camilleri also identifies some of the key elements required for responsible governance of artificial intelligence (Figure 3). When properly implemented, such frameworks can help align business practices with ethical principles and foster greater trust among users and the public.

However, several limitations undermine the effectiveness of corporate governance in practice. One significant issue is the reliance on self-regulation. Many organizations design and enforce their own ethical guidelines without external oversight, which can create conflicts of interest. Companies may be incentivized to prioritize business objectives over ethical considerations, particularly when governance mechanisms lack independence or authority. Hagendorff challenges this approach by arguing that voluntary ethics guidelines often fail to produce meaningful behavioural change, as they do not impose binding constraints or consequences for breaking said guidelines (Hagendorff 2020).

Another concern is the phenomenon of “ethics washing”, where organizations publicly promote their commitment to ethical AI without implementing substantive changes to their practices. In such cases, governance frameworks may serve primarily as a reputational tool, rather than showing accountability. This can create a false sense of assurance among stakeholders while allowing problematic practices to persist. The lack of transparency surrounding internal decision making processes complicates efforts to assess whether organizations are genuinely adhering to ethical standards (Schultz, Conti & Seele 2025).

Accountability in AI systems is also inherently complex due to the distributed nature of decision making within organizations. AI development often involves multiple teams, including data scientists, engineers, managers, and executives, each contributing to different stages of the system lifecycle. This diffusion of responsibility can make it difficult to assign accountability when ethical issues arise, particularly in cases involving unintended or biased outcomes.

Without clear lines of responsibility, organizations may struggle to ensure that ethical standards are consistently upheld.

Overall, while corporate governance frameworks represent an important step toward responsible AI, they are often insufficient on their own to address the underlying tension between business incentives and ethical obligations. Strengthening accountability will require not only improved internal governance but also greater external oversight and auditing, clearer regulatory standards, and increased transparency in organizational practices.

3.5 AI in Government and Defence Contracts in the US

The deployment of artificial intelligence within government and defence contexts in the US represents one of the most complex and ethically sensitive applications. Increasingly, governments rely on private sector companies to develop AI systems for national security, intelligence analysis, surveillance, and military operations. While these partnerships can enhance strategic capabilities and operational efficiency, they also intensify the tension between business incentives and ethical responsibility, especially because of the high-stakes nature of defence applications.

A central issue in this context is the question of accountability. When private companies develop AI systems for government use, responsibility for the outcomes of those systems can become fragmented. Tillipman highlights that while companies may retain certain legal rights in government contracts, they often have limited control over how their technologies are ultimately deployed (Tillipman 2026). This creates a situation in which harmful or unethical outcomes: such misuse of these tools, or errors in automated decision-making may not be clearly attributable to a single actor. The diffusion of responsibility between developers and government users complicates efforts to enforce ethical standards and ensure accountability.

Another significant concern is the impact of AI on human decision-making in military contexts. Traditional frameworks of military ethics emphasize human judgement, moral reasoning, and adherence to legal principles such as proportionality and distinction. However, the integration of AI into decision support systems and autonomous operations raises the risk that the human aspect of moral decision-making be decreased or even outright eliminated. Johnson argues that reliance on AI during military crises could lead to a form of “moral distancing”, where human actors become less directly engaged with the ethical consequences of their decisions. This raises important questions about whether AI systems should be permitted to influence, or even make, decisions involving the use of force.

The emergence of advanced AI technologies, including generative and agentic systems, further amplifies these concerns. In national security settings, such systems may be used for intelligence analysis and strategic planning, offering significant advantages in speed and scale. However, as noted by Joshi, these capabilities also introduce risks related to unpredictability, system autonomy, and the potential for unintended escalation of conflict (Joshi 2025). Similarly, research on AI in cyber and kinetic warfare highlights the possibility that automated systems could accelerate decision-making processes beyond human control, increasing the likelihood of errors (D'Arcy 2025).

Efforts to mitigate these risks have included the adoption of technical safeguards such as AI “red-teaming”, where systems are tested for vulnerabilities, biases, and potential misuse before deployment. While such practices can identify certain risks, they are not a comprehensive solution. Feffer et al. argue that red-teaming may, in some cases, function as a

form of "security theatre", providing a perception of safety without addressing deeper structural and ethical issues. As a result, reliance on technical testing alone is insufficient to ensure responsible AI deployment in defence contexts (Feffer et al. 2024).

Metric	*Purpose*	*AI-DSS technical features*	*Benchmark/threshold*
Decision-making accountability	Ensures human operators retain authorship and responsibility for decisions and outcomes	Transparent logs, justification templates, real-time AI tracking, time-stamped actions, and override ratios	
Trust in DSS calibration	Assesses whether humans exhibit appropriate trust in AI outputs, avoiding overreliance	Confidence and uncertainty indicators, explainability features, mandatory override justifications, and user trust metrics	Real-time decision tracking. E.g. an override rate >25% indicates healthy human engagement
Transparency and understandability	Assesses whether the system presents its reasoning and confidence clearly enough to support timely, ethically informed human decisions under pressure	Decision paths, interactive explanations with layered detail, and tracking AI confidence–human acceptance correlations	Confidence-scoring system (0–100%)
Time for deliberation	Evaluates how effectively the system balances moral reflection with time-sensitive decision demands	Integrated workflows, streamlined interfaces, and prioritized alerts for timely interaction	
Ethical override frequency	Measures how often humans override AI recommendations	Override mechanisms with clear feedback loops and justifications for AI recommendations	Override pattern analysis e.g. 70–90% acceptance rate for high confidence
Identify and mitigate human cognitive bias[a]	Evaluates the system's ability to reduce human biases such as automation bias and anthropomorphism	Counterfactual scenarios, alternative options, training, bias alerts	
Collaborative feedback loops	Monitors the quality and frequency of human–AI interaction in decision refinement	Feedback prompts, team dashboards, and adaptive, user-informed updates	
Psychological impact on users	Tracks users' moral stress, emotional well-being, and satisfaction during interaction	Surveys, behaviour monitoring, and stress-reducing interface design	
Ethical adaptability in multi-agent scenarios	Tests the system's ability to dynamically adjust to context-specific dilemmas or feedback	Real-time updates, contextual learning, and situation-specific adaptation	

[a] Bias in decision-making may arise from the decision-maker, the surrounding context or a misalignment between the two. In both experimental settings and operational environments, decisions are frequently made under conditions of novelty, ambiguity and competing objectives. See Baruch Fischhoff, *Prejudices about bias*, working paper no. 85–3 (Eugene, OR: Decision Research, 1985), http://doi.org/10.2139/ssrn.1560126.

FIGURE 4 METRICS AND TECHNICAL FEATURES FOR ENHANCING HUMAN MORAL AGENCY IN HUMAN–AI DSS INTERACTION.(JOHNSON 2026)

Overall, the use of AI in government and defence contracts in the US illustrates an intensified form of the broader tension between business incentives and ethical responsibility. Johnson outlines key metrics and technical features designed to enhance human moral agency in AI-assisted military decision-making in Figure 4 (Johnson 2026). The combination xgof high financial risk, national security priorities, and limited transparency creates an environment in which ethical considerations may be subordinated to strategic objectives. Addressing these challenges requires stronger governance frameworks, clearer accountability mechanisms, and sustained emphasis on maintaining human oversight in critical decision-making processes.

3.6 Effectiveness of Current AI Ethics Frameworks

AI ethics frameworks have emerged as a central tool for promoting responsible development and deployment of artificial intelligence systems. Governments, international organizations, and industry bodies have introduced many guidelines aimed at ensuring that AI technologies align with societal values such as fairness, transparency, accountability, and human oversight. For example, Australia's AI Ethics Principles provide a comprehensive set of voluntary guidelines intended to guide organizations in the ethical design and use of AI. While such

frameworks play an important role in shaping discourse and raising awareness, their effectiveness in practice remains limited.

One of the primary strengths of AI ethics frameworks is their ability to establish a shared set of normative expectations. By articulating key principles, these frameworks help organizations recognize potential ethical risks and encourage the adoption of responsible practices. They also contribute to the development of a common language around AI ethics, which facilitates collaboration between policymakers, industry, and academic research. In this sense, ethics frameworks serve as an important foundation for governance, particularly in a rapidly evolving technological landscape.

However, significant limitations reduce the practical impact of these frameworks. A major issue is their lack of enforceability. Most AI ethics guidelines, including Australia's, are voluntary and do not impose binding legal obligations on organizations and businesses. As a result, compliance is largely dependent on organizational goodwill, which may be insufficient in the face of strong commercial incentives. Hagendorff argues that the proliferation of AI ethics principles has not translated into meaningful behavioural change, as many guidelines remain abstract and lack mechanisms for accountability or enforcement (Hagendorff 2020). This gap between principle and practice inhibits the ability of such frameworks to govern AI development.

Another limitation is the inherent vagueness of many ethical principles Concepts such as fairness, transparency, and accountability are open to interpretation and may be implemented in inconsistent ways across organizations. This ambiguity allows companies to selectively adopt aspects of ethical frameworks that align with their business objectives while neglecting more challenging requirements. In some cases, organizations may engage in "ethics washing", publicly endorsing ethical principles without making substantive operational changes, therefore reducing the credibility of such frameworks.

Furthermore, existing ethics frameworks often struggle to keep pace with emerging AI technologies, particularly in high-risk domains such as generative AI, autonomous systems, and defence applications. As discussed in previous sections, these technologies introduce new challenges related to autonomy, scalability, and unpredictability that are not always adequately addressed by current guidelines. This raises questions about whether existing frameworks are sufficiently adaptable to govern increasingly complex AI systems.

Overall, while AI ethics frameworks are a necessary component of responsible AI governance, they are not sufficient on their own. Their effectiveness is constrained by their voluntary nature, lack of enforcement mechanisms, and inability to fully address the complexities of modern AI systems. Strengthening their impact will require integration with regulatory measures, organizational accountability structures, and ongoing adaptation to technological developments.

4. Conclusion and Recommendations

This report has examined the tension between business incentives and ethical responsibilities in the development and deployment of artificial intelligence systems. The findings demonstrate that while AI offers significant economic and operational benefits, its adoption is accompanied by complex ethical challenges that are often difficult to address in practice. Organizations operate within competitive environments that prioritize efficiency, innovation, and profitability,

which can conflict with ethical principles such as fairness, transparency, and accountability. As a result, ethical considerations are frequently deprioritized, especially when they impose additional costs or delay deployment.

The analysis has shown that the challenges of implementing responsible and ethical AI are both technical and organizational in nature. Ambiguity in ethical principles, resource constraints, and limitations in current AI technologies all contribute to the difficulty of translating ethical guidelines into practice. Furthermore, companies may engage in unethical AI practices as a rational response to market pressures and weak regulatory enforcement. Although corporate governance mechanisms and ethics frameworks provide some level of oversight, they are often insufficient due to issues such as self-regulation, lack of accountability, and the risk of ethics washing. These challenges are further exacerbated in high-stakes contexts such as government and defence contracts, where the consequences of unethical AI deployment can be particularly severe.

From the findings, some recommendations can be made. First, there is a need for stronger regulatory frameworks that move beyond voluntary guidelines and establish clear, enforceable standards for ethical AI. Second, Organizations should integrate ethical considerations into their core business strategies rather than treating them as external or secondary concerns. This includes accountability mechanisms, performance metrics, and leadership responsibility. Third, greater transparency is required in AI development and deployment processes to enable external scrutiny and build public trust. Finally, ongoing investment in research and development is necessary to address technical challenges such as bias, explainability, and system reliability, particularly in emerging areas like generative and autonomous AI.

In conclusion, achieving ethical and responsible AI requires a coordinated effort across industry, government, and academic research. While existing frameworks provide a valuable starting point, meaningful progress will depend on aligning business incentives with ethics through stronger governance, regulation and organizational changes.

5. Acknowledgements/Resources

I used Google's Scholar Labs tool to help me find relevant papers for the research question 5 in the COS30019 Assignment 1 outline pdf.

References

Australia's AI Ethics Principles, 2025, Australian Government, viewed 3/15/2026, <https://www.industry.gov.au>.
Camilleri, MA 2024, 'Artificial intelligence governance: Ethical considerations and implications for social responsibility', *Expert systems*, vol. 41, no. 7, p. e13406.
D'Arcy, AE 2025, 'From Code to Combat: Introducing the Role of AI in Cyber and Kinetic Warfare in 2024,' *2025 Cyber Research Conference-Ireland (Cyber-RCI)*, IEEE, 1-8.
Feffer, M, Sinha, A, Lipton, ZC & Heidari, H 2024, 'Red-Teaming for Generative AI: Silver Bullet or Security Theater? arXiv', *Preprint posted online January*, vol. 29,
Hagendorff, T 2020, 'The ethics of AI ethics: An evaluation of guidelines', *Minds and machines*, vol. 30, no. 1, pp. 99-120.
Johnson, J 2026, 'Can AI behave ethically during military crises? Preserving human moral agency', *International Affairs*, vol. 102, no. 1, pp. 63-83.
Joshi, S 2025, 'Agentic Generative AI in US National Security and Defense: Tools, Frameworks, and Policy Recommendations', *Authorea Preprints*,
Maslej, N, Fattorini, L, Perrault, R, Gil, Y, Parli, V, Kariuki, N, Capstick, E, Reuel, A, Brynjolfsson, E & Etchemendy, J 2025, 'Artificial intelligence index report 2025', *arXiv preprint arXiv:2504.07139*,
Méndez-Suárez, M, de Obesso, MdlM, Márquez, OC & Palacios, CM 2023, 'Why do companies employ prohibited unethical artificial intelligence practices?', *IEEE Transactions on Engineering Management*, vol. 71, pp. 12218-12227.
Purwanto, ANI, Fauzan, M, Widya, T & Azzaky, NS 2024, 'Ethical Implications and challenges of AI implementation in business operations', *TechComp Innovations: Journal of Computer Science and Technology*, vol. 1, no. 2, pp. 68-82.
Schultz, MD, Conti, LG & Seele, P 2025, 'Digital ethicswashing: A systematic review and a process-perception-outcome framework', *AI and Ethics*, vol. 5, no. 2, pp. 805-818.
Tillipman, J 2026, 'What Rights Do AI Companies Have in Government Contracts?',
Wright, SA & Schultz, AE 2018, 'The rising tide of artificial intelligence and business automation: Developing an ethical framework', *Business Horizons*, vol. 61, no. 6, pp. 823-832.